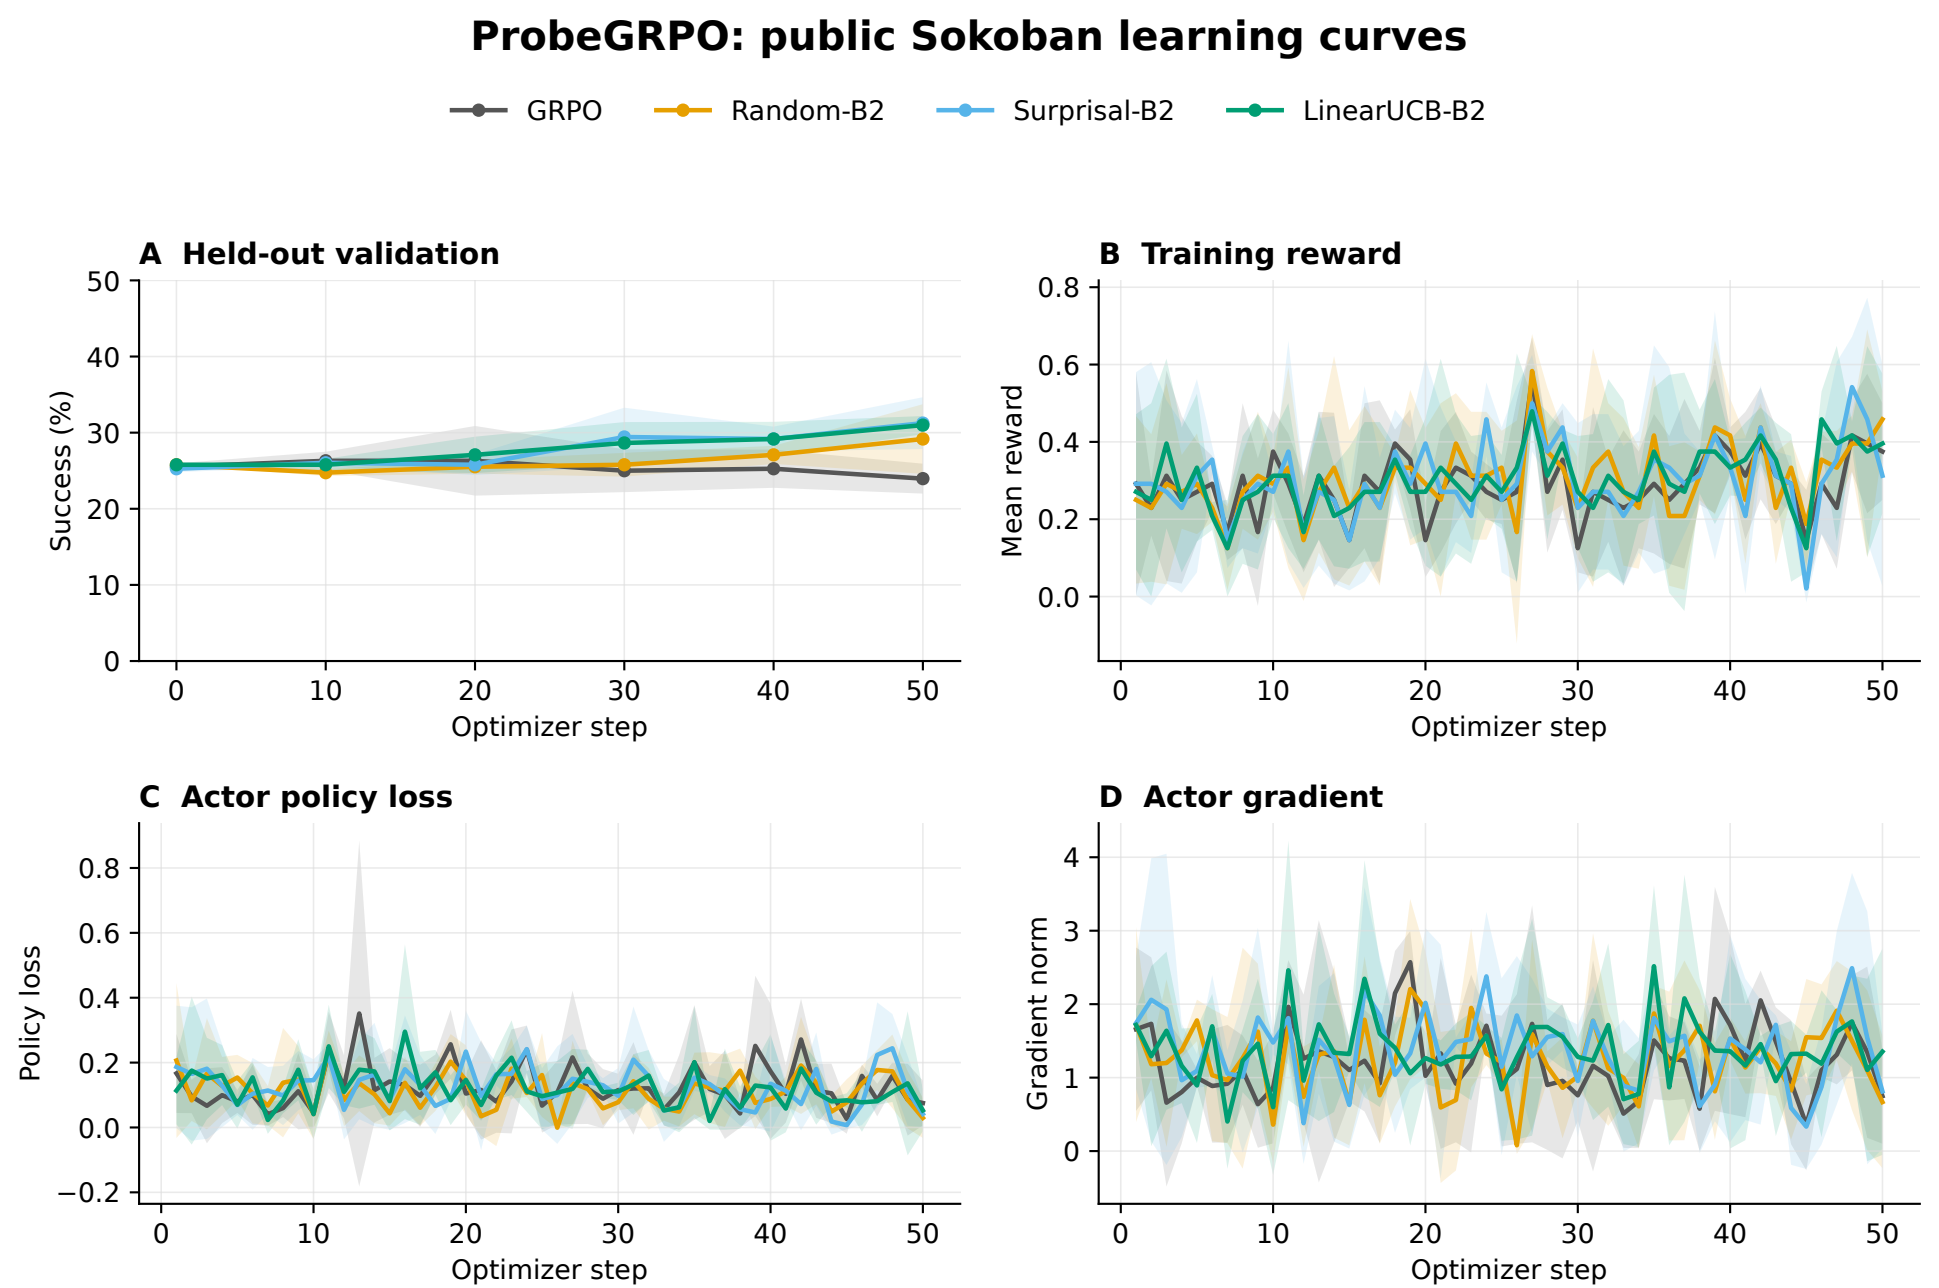
ProbeGRPO: public Sokoban learning curves
GRPO
Random-B2
Surprisal-B2
LinearUCB-B2
A Held-out validation
Success (%)
Optimizer step
B Training reward
Mean reward
Optimizer step
C Actor policy loss
Policy loss
Optimizer step
D Actor gradient
Gradient norm
Optimizer step
Means across seeds 17, 42, 101; shading = sample SD (n=3); no smoothing.